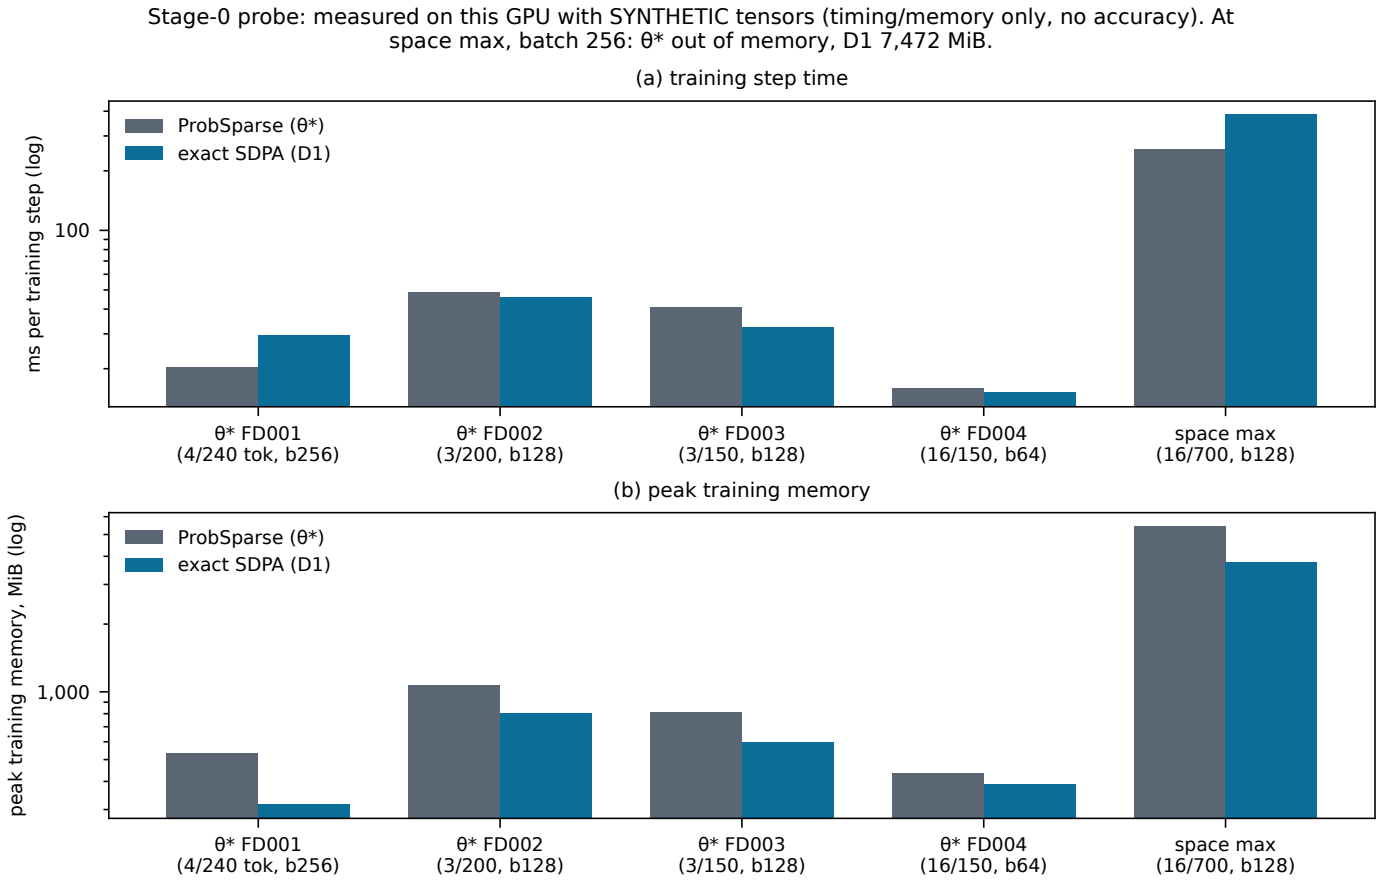
Stage-0 probe: measured on this GPU with SYNTHETIC tensors (timing/memory only, no accuracy). At space max, batch 256: θ* out of memory, D1 7,472 MiB.
(a) training step time
ms per training step (log)
100
ProbSparse (θ*)
exact SDPA (D1)
θ* FD001 (4/240 tok, b256)
θ* FD002 (3/200, b128)
θ* FD003 (3/150, b128)
θ* FD004 (16/150, b64)
space max (16/700, b128)
(b) peak training memory
peak training memory, MiB (log)
1,000
ProbSparse (θ*)
exact SDPA (D1)
θ* FD001 (4/240 tok, b256)
θ* FD002 (3/200, b128)
θ* FD003 (3/150, b128)
θ* FD004 (16/150, b64)
space max (16/700, b128)